

Privacy Enhancing Technologies

1 Commitment Schemes

Def. Commitment scheme algorithm $\text{Commit} : \mathcal{M} \times \mathcal{R} \rightarrow \mathcal{C}$

- **(Statistically) Hiding:** c leaks nothing about m
- **(Computationally) Binding:** No PTT adversary can produce $c_1 = c_2$ for $m_1 \neq m_2$

Random oracle model Replace hash function by *completely random function* $H : \mathcal{X} \rightarrow \mathcal{Y}$.

Thm. $\text{Commit}(m, r) := H(m, r)$ is stat. hiding and comp. binding in ROM ($|\mathcal{K}|$ and $|\mathcal{C}|$ have size $2^{\text{poly}(\lambda)}$)

Def. Pedersen Comm. Let \mathbb{G} be group of prime order q , let $g, h \in \mathbb{G}$ generators s.t. $g^x = h$ and x unknown.

$$\text{Commit}(m, r) := g^m h^r = g^{m+r} \cdot x$$

Properties

- Linearly homomorphic $(\text{Commit}(m_0, r_0) \cdot \text{Commit}(m_1, r_1)) = \text{Commit}(m_0 + m_1, r_0 + r_1)$
- Statistically hiding (for any c and m there is a unique r)
- Computationally binding (same advantage as in DLog)

How to **break DLog**: find two representations of same element: $g^{m_0} h^{r_0} = c = g^{m_1} h^{r_1}, m_0 \neq m_1 \implies m_0 + x \cdot r_0 = m_1 + x \cdot r_1 \implies x = \frac{m_1 - m_0}{r_0 - r_1} \bmod q$

Def. DDH Problem Consider the distributions

$$D_0 = \{(g^\alpha, g^\beta, g^{\alpha\beta}) : \alpha, \beta \leftarrow \mathbb{Z}_q\}$$

$$D_1 = \{(g^\alpha, g^\beta, g^\gamma) : \alpha, \beta, \gamma \leftarrow \mathbb{Z}_q\}$$

$\implies \text{AdvDDH}(\mathcal{A}, \mathbb{G}) = |\Pr[W_0] - \Pr[W_1]|$, where W_0, W_1 are events that \mathcal{A} outputs 1 given a sample from D_0 or D_1 .

DDH Assumption $\text{AdvDDH}(\mathcal{A}, \mathbb{G}) = \text{negl}(\lambda)$ for all PTT \mathcal{A}

Thm. If DDH Assumption holds and $H : \mathbb{Z}_q \times \mathcal{X} \rightarrow \mathbb{G}$ is a random function, then **PRF** $F(k, x) = H(x)^k$ is secure.

Vector commitment Produces commitment c of size λ and opening proof π of size $= (\log n \cdot \lambda)$: $H(m_1), \dots, H(m_n)$ are leaves of binary tree. Each node with children v_l, v_r has value $H(v_l, v_r)$. Commitment c is root hash.

To open at position i , walk from leaf i to root, revealing values of the sibling nodes along the path.

2 Zero Knowledge Proofs

NP Language NP is class of languages \mathcal{L} s.t. there exists efficient, deterministic algo M s.t.

$$x \in \mathcal{L} \iff \exists w \in \{0, 1\}^{\text{poly}(|x|)} \text{ s.t. } M(x, w) = 1$$

Def. Interactive Proof unbounded prover, efficient verifier

- **Complete:** $\forall x \in \mathcal{L} : \Pr[\langle P, V \rangle(x) = 1] \geq \frac{2}{3}$
- **Sound:** $\forall x \notin \mathcal{L}, \forall$ provers $P^* : \Pr[\langle P^*, V \rangle(x) = 1] \leq \frac{1}{3}$

Def. ZK Proof Interactive proof system $\langle P, V \rangle$ that satisfies completeness and soundness, and \forall (malicious) V^* :

\exists simulator Sim_{V^*} s.t. $\forall x \in \mathcal{L} : \text{View}[\langle P, V^* \rangle] \approx^c \text{Sim}_{V^*}(x)$

Ex. ZK proof for Ham-Cycle Assuming comm. scheme w/ perfect binding, computational hiding: valid ZK proof w/ perfect completeness and soundness error $\leq \frac{1}{2}$

1. $P(G, \text{Ham-Cycle})$: pick random permutation, insert edges into permuted adjacency matrix and commit to each value
2. $V(G)$: sample random challenge $c \leftarrow \mathbb{S} \{0, 1\}$
3. P : if $c = 0$, open commitments. Else, open all commitments where edges are in hamiltonian cycle
4. V : if $c = 0$, check if graphs are isomorphic. Else, check that a cycle was opened

Sim(G):

1. Guess $\hat{c} \leftarrow \mathbb{S} \{0, 1\}$. If $\hat{c} = 0$, commit to random permutation of G . Else, commit to complete graph
2. Send commitments and receive c
3. If $c \neq \hat{c}$ abort and restart
4. Else, open requested commitments and output the view

Parallel vs. in sequence In sequence: run the protocol λ times one-after-the-other. Parallel: run λ instances with λ provers interacting with their verifiers. Tricky to build an efficient simulator for parallel, since if verifier is malicious, it can sample the challenges in a way unknown to the verifier and easily detect a simulated view.

Def. Proof of Knowledge Proof system $\langle P, V \rangle$ for NP lang \mathcal{L} is proof of knowledge w/ knowledge error ϵ if efficient extractor Ext exists s.t. $\forall x, P^*$:

$$\Pr[M(x, w) = 1 : w \leftarrow \text{Ext}^{P^*}(x)] \geq \Pr[\langle P, V \rangle(x) = 1] - \epsilon$$

The extractor can rewind the prover.

Def. Sigma protocol Three-round protocol $\langle P, V \rangle$ of the form *commitment* $u \rightarrow \text{challenge } c \rightarrow \text{response } z$ with:

- perfect completeness: accepts all valid statements
- special soundness: there exists Ext , s.t. given 2 accepting transcripts $(u, c, z), (u, c', z')$ outputs w s.t. $M(x, w) = 1$
- honest verif. ZK: Sim outputs computationally indistinguishable transcript from a honest verifier's view

Thm. Sigma prot. gives interact. proof w/ sound. err. $\leq \frac{1}{|C|}$

Non-interactive NIZK Simulator can program the RO for verifier's RO queries, as long as responses are indist. from random

Fiat-Shamir transform Turns Sigma-protocol into NIZK: compute challenge as hash of prior messages of prover ($c = H(x, u)$)

For NIZK, we require **negl. soundness and completeness!**

3 SNARGs

Lem. Equality Test To check if polynomials f and g are equal, check equality at a random point. This has perfect completeness and soundness error $\frac{d}{|\mathbb{F}|}$.

Fundamental theorem of algebra Any non-zero polynomial over field \mathbb{F} of degree d has at most d roots over \mathbb{F} .

Lem. Schwartz-Zippel Let $f(X_1, \dots, X_n)$ be multivariate polynomial of total degree d . Let Ω a finite subset of \mathbb{F} and $r_1, \dots, r_n \in \Omega$ uniformly random. Then $\Pr[f(r_1, \dots, r_n) = 0] \leq \frac{d}{|\Omega|}$

Thm. Let $\Omega \subseteq \mathbb{F}$, $|\Omega| = k \leq d$. Let $Z_\Omega(X) = \prod_{a \in \Omega} (X - a)$. Then: f is zero on $\Omega \iff f(X)$ divisible by $Z_\Omega(X)$

Thm. PCP Let L be an NP language. There exist two efficient algorithms $(\mathcal{P}, \mathcal{V})$ defined as follows:

- Prover \mathcal{P} is deterministic algo that takes as input a statement $x \in \{0, 1\}^n$, a witness $w \in \{0, 1\}^h$ and outputs a bitstring $\pi \in \{0, 1\}^m$, where $h, m = \text{poly}(n)$. We refer to π as the proof string.

- Verifier algo \mathcal{V}^π is randomized and takes as input a statement $x \in \{0, 1\}^n$ and has oracle access to proof string $\pi \in \{0, 1\}^m$. Verifier reads $O(1)$ bits of π (chosen non-adaptively).

- **Complete:** $\forall x \in \mathcal{L}$ if w witness for x : $\Pr[\mathcal{V}^\pi(x) = 1 : \pi \leftarrow \mathcal{P}(x, w)] = 1$

- **Sound:** if $x \notin \mathcal{L}$, then for all $\pi \in \{0, 1\}^m$: $\Pr[\mathcal{V}^\pi(x) = 1] \leq \frac{1}{2}$

Prove that $f(z) = y$ at some point z using ZeroTest: we can write $f(X) - y = (X - z) \cdot q(X)$, i.e. $q(X) = \frac{f(X) - y}{X - z}$

PLONK IOP Represent an arithmetic circuit as polynomials, using $w \in \mathbb{F}$ a root of unity (i.e. $w^{3|C|+|I|} = 1$). When applying Fiat-Shamir, this becomes a SNARG for NP.

- $f_{trace}(w^{-j})$ = input j
- $f_{trace}(w^{3l})$ = left input to gate l
- $f_{trace}(w^{3l+1})$ = right input to gate l
- $f_{trace}(w^{3l+2})$ = output of gate l

To verify, then check the following:

- Output of last gate is 0 ($f_{trace}(w^{3|C|-1}) = 0$)
- Public inputs: $f_{trace} - f_{in} = 0$ (ZeroTest)
- each gate is evaluated correctly (ZeroTest, using selector polynomial): $f_{sel}(x) \cdot (f_{trace}(x) + f_{trace}(wx)) + (1 - f_{sel}(x)) \cdot (f_{trace}(x) \times f_{trace}(wx)) = f_{trace}(w^2 x)$
- wiring is implemented correctly (outputs = inputs)

4 Private Information Retrieval (PIR)


2-server PIR View the DB as a matrix X of $\sqrt{n} \times \sqrt{n}$ bits. Then, using $O(n)$ communication, load the j -th row and read the i -th bit to retrieve X_{ij} .

- **Query**((i, j): $q_0 \leftarrow \mathbb{Z}_2^{\sqrt{n}}, q_1 = q_0 \oplus e_j$
- **Answer**(X, q_b): $a_b \leftarrow X \cdot q_b \in \mathbb{Z}_2^{\sqrt{n}}$
- **Reconstruct**(a_0, a_1): $x_i \leftarrow (a_0 \oplus a_1)_i$


Linearly homomorphic encryption

$$\text{Enc}(k, m_1) + \text{Enc}(k, m_2) = \text{Enc}(k, m_1 + m_2)$$

Single-server PIR Using LHE and load balancing:



Recursive PIR View X as $n^{2/3} \times n^{1/3}$ matrix and recurse:



Watch out for ciphertext expansion!

5 Oblivious RAM

Def. Privately read and write to private RAM

- **Correctness:** ORAM answers correctly if RAM is honest
- **Security:** For any two poly-sized seq. of equal length: $\{\text{Access}(op_1), \dots, \text{Acc}(op_k)\} \approx^c \{\text{Access}(op'_1), \dots, \text{Acc}(op'_k)\}$ i.e. only the number of operations is leaked

A \sqrt{n} Construction Any $k \leq n$ ops that access distinct locations are indistinguishable. The client keeps stash of \sqrt{n} words in internal state.

while (true): The RAM is shuffled according to some random permutation π (oblivious sorting, data-independent). If some value is modified in the stash, update the RAM. Process \sqrt{n} ops:

- if address in stash, read random address and discard
- else, read address and store in stash
- if op is write, update in stash

Amortized cost per access: $O(\sqrt{n}/k)$ where k is size of stash and $k = \sqrt{n}$ for balanced costs. Shuffle in $O(n \log^2 n)$ time, then process \sqrt{n} ops with 1 RAM access each.

Lower bound $O(\log n)$ accesses per op, there are schemes that achieve this bound. For $O(\text{polylog} n)$ use tree-based ORAM.

6 Multi-party computation (MPC)

Def. n Parties P_1, \dots, P_n with inputs x_1, \dots, x_n want to compute $y \leftarrow f(x_1, \dots, x_n)$.

Adv should learn nothing more about honest parties' inputs than if they were interacting with trusted third party. An adversary learns 1) inputs of corrupted parties and 2) output y of func f .

Def. Semi-honest security Adversary follows protocol spec, but can do additional poly-time local computation. Tries to learn honest parties' inputs. The view of the corrupted parties C should be indistinguishable from simulation.

$$\text{Sim}(C, \underbrace{\{x_i : i \in C\}}_{\text{the inputs of corrupted parties}}, \underbrace{y = f(x_1, \dots, x_n)}_{\text{the output of the computation}}) \approx \underbrace{\{\text{View}_i : i \in C\}}_{\text{the views of all corrupted parties in a real protocol execution}}$$

Def. Malicious security Adversary may arbitrarily deviate from protocol to learn about honest parties' inputs or fool them into accepting incorrect output.

7 Examples

New DDH triples Build new DDH triple from (X, Y, Z) : sample $r \leftarrow \mathbb{Z}_q$, set $X' = X \cdot g^r, Z' = Z \cdot Y^r$

PRF security Using PRF $F(k, x) = H(x)^k$, PRF adversary has access to RO for H . DDH adversary simulates a PRF challenger and receives (X, Y, Z) . When getting RO query m_i , sample random $r_i \leftarrow \mathbb{Z}_q$, set $H(m_i) = X \cdot g^{r_i}$ and store (m_i, r_i) .

When getting a PRF query for m_j , return $Z \cdot Y^{r_j}$. Simulation is perfect, since RO responses are random. If $b = 0$, Z is random and so is $Z' = Z \cdot Y^{r_j}$.

If $b = 1$, $Z' = (X')^\beta = H(m_j)^\beta = F(\beta, m_j)$.

SNARGs Prove that $\text{fib}(100) = x$. Prover puts the following polynomials in the box: $f_{\text{fib}}(X)$ and $q(X) = \frac{f_{\text{fib}}(X+2) - f_{\text{fib}}(X) - f_{\text{fib}}(X+1)}{\prod_{i=0}^{98}(X-i)}$. The verifier queries $f_{\text{fib}}(0), f_{\text{fib}}(1), f_{\text{fib}}(100)$ and $f_{\text{fib}}(r), f_{\text{fib}}(r+1), f_{\text{fib}}(r+2), q(r)$ and checks $f_{\text{fib}}(r+2) - f_{\text{fib}}(r) - f_{\text{fib}}(r+1) = q(r) \cdot \prod_{i=0}^{98}(r-i)$. This has perfect completeness and soundness error $\frac{2d}{|\mathbb{F}|}$.

PIR from FHE Given secret-key FHE scheme that supports enc, NOT, AND, OR: build 1-server PIR for DB of n bits.

Write $i \in [n]$ as binary string $i = i_1 i_2 \dots i_{\log n}$ and send to server, along with enc. of 0 and 1.

For each index $j = j_1 j_2 \dots j_{\log n}$ the server replaces j_i by corresponding ciphertext and homomorph. evaluates $C_{EQ}(i, j)$ over encrypted bits. Then, multiply this with ciphertext of bit X_j and sum up all ciphertexts (with OR).

PIR from DPF A DPF creates keyed functions $\text{Eval}(k_0, \cdot)$, $\text{Eval}(k_1, \cdot)$ s.t. $\text{Eval}(k_0, \cdot) + \text{Eval}(k_1, \cdot) = P_{x,y}(\cdot)$.

To retrieve i -th bit, generate keys $(k_0, k_1) \leftarrow \text{Gen}(i, 1)$.

$$(\sum_j \text{Eval}(k_0, j) \cdot X_j) \oplus (\sum_j \text{Eval}(k_1, j) \cdot X_j) = \sum_j P_{i,1}(j) \cdot X_j$$

For *PIR with keywords*, generate keys $(k_0, k_1) \leftarrow \text{Gen}(w', 1)$.